

Importing libraries

In [1]:

```
import numpy as np, pandas as pd
import matplotlib.pyplot as plt, seaborn as sns

import warnings

from sklearn.model_selection import train_test_split

from sklearn.preprocessing import StandardScaler

from sklearn.feature_selection import RFE
from sklearn.linear_model import LinearRegression
from statsmodels.stats.outliers_influence import variance_inflation_factor
import statsmodels.api as sm

from sklearn.metrics import r2_score
```

In [2]:

```
warnings.filterwarnings('ignore')
```

Import DATASET

In [3]:

```
df = pd.read_csv("AutoData (1).csv")
```

In [4]:

```
# Viewing an overview of data
```

```
df.head()
```

Out[4]:

	symboling	make	fueltype	aspiration	doornumber	carbody	drivewh
0	3	alfa-romero giulia	gas	std	two	convertible	r
1	3	alfa-romero stelvio	gas	std	two	convertible	r
2	1	alfa-romero Quadrifoglio	gas	std	two	hatchback	r
3	2	audi 100 ls	gas	std	four	sedan	f
4	2	audi 100ls	gas	std	four	sedan	4

5 rows × 25 columns



Basic EDA

In [5]:

```
# Checking the data-types of columns (checking for data-type mismatch)
```

```
df.dtypes
```

Out[5]:

```
symboling          int64
make              object
fueltype          object
aspiration        object
doornumber        object
carbody           object
drivewheel        object
enginelocation    object
wheelbase         float64
carlength         float64
carwidth          float64
carheight         float64
curbweight        int64
enginetype        object
cylindernumber    object
enginesize        int64
fuelsystem        object
boreratio         float64
stroke            float64
compressionratio  float64
horsepower        int64
peakrpm           int64
citympg           int64
highwaympg        int64
price             float64
dtype: object
```

In [6]:

```
# Dimensions of dataset
```

```
df.shape
```

Out[6]:

```
(205, 25)
```

Since the data-set is quite small, any redundant/high class imbalanced columns must be handled smartly and not dropped

In [7]:

```
# df.info()
```


In [8]:

```
# Basic statistical description of numeric columns
```

```
df.describe().round(2)
```

Out[8]:

	symboling	wheelbase	carlength	carwidth	carheight	curbweight	engine
count	205.00	205.00	205.00	205.00	205.00	205.00	205.00
mean	0.83	98.76	174.05	65.91	53.72	2555.57	120.00
std	1.25	6.02	12.34	2.15	2.44	520.68	40.00
min	-2.00	86.60	141.10	60.30	47.80	1488.00	60.00
25%	0.00	94.50	166.30	64.10	52.00	2145.00	90.00
50%	1.00	97.00	173.20	65.50	54.10	2414.00	120.00
75%	2.00	102.40	183.10	66.90	55.50	2935.00	140.00
max	3.00	120.90	208.10	72.30	59.80	4066.00	320.00



In [9]:

```
# Finding the count and an overview of all the unique segments present in categ
```

```
for i in df.select_dtypes('object').columns:
    print(i, '\n')
    print(df[i].nunique())
    print(df[i].unique()[:5])
    print('\n\n')
```

make

147

```
['alfa-romero giulia' 'alfa-romero stelvio' 'alfa-romero Quadrif
oglio'
 'audi 100 ls' 'audi 100ls']
```

fueltype

2

```
['gas' 'diesel']
```

aspiration

2

```
['std' 'turbo']
```

doornumber

2

```
['two' 'four']
```

carbody

5

```
['convertible' 'hatchback' 'sedan' 'wagon' 'hardtop']
```

drivewheel

3

```
['rwd' 'fwd' '4wd']
```

enginelocation

```
2  
['front' 'rear']
```

enginetype

```
7  
['dohc' 'ohcv' 'ohc' 'l' 'rotor']
```

cylindernumber

```
7  
['four' 'six' 'five' 'three' 'twelve']
```

fuelsystem

```
8  
['mpfi' '2bbl' 'mfi' '1bbl' 'spfi']
```



Visualization



Uni-variant analysis



In [10]:




```
# Checking distribution for numeric columns
```



```
for i in df.select_dtypes('number'):
    sns.distplot(df[i])
    plt.show()
```



















Risk factor (symboling) in our data is quite risky as curve is higher at +ve side

- There are no records at -3 i.e Safest side
- Also, this column can be treated as a categorical one

Bore Ratio = Bore (Diameter of cylinder)/Stroke (Length of Piston)

- So expect bore ratio and stroke to have high co-relation (negative), if so stroke is a redundant column and can be dropped

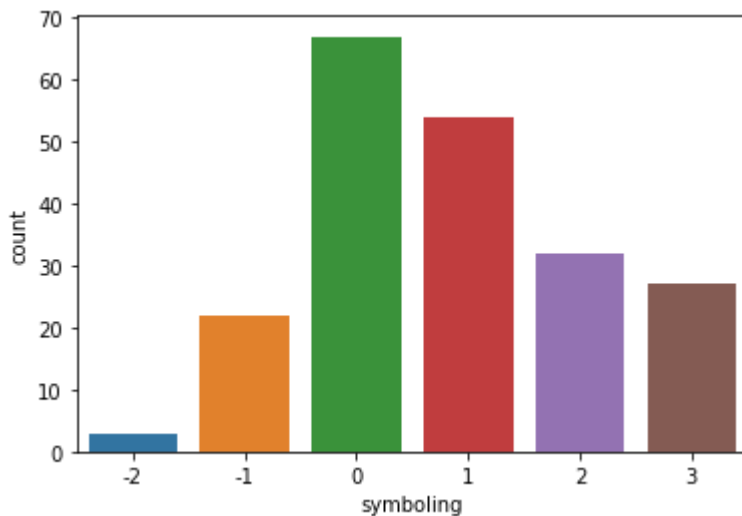
Categorical Analysis

Not analysing for make as seen earlier, it has too many columns not easy analysis can be performed

In [11]:

```
# Checking the share of Fuel-type
```

```
sns.countplot(x = 'symboling', data = df)  
plt.show()
```




Many risky cars rated vehicles are present in data

In [12]:

```
# Checking the share of Fuel-type
```

```
sns.countplot(x = 'fueltype', data = df)  
plt.show()
```

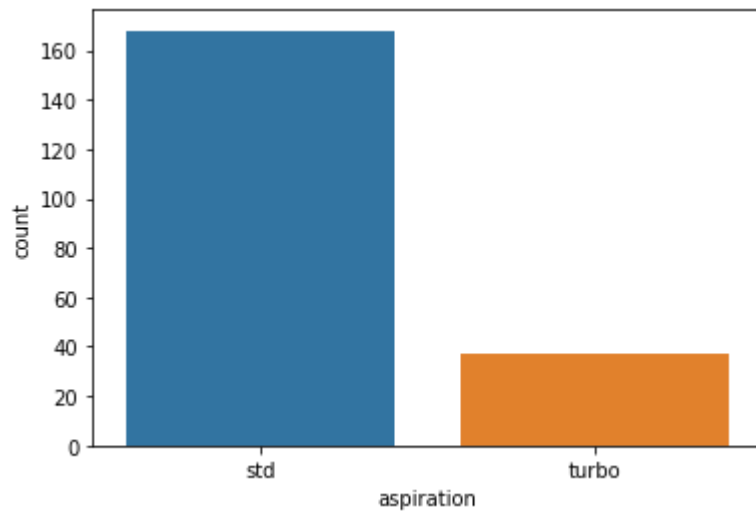


Majority of the data present is of 'gas' type, so prices will be skewed towards it, hence mean (average) can be a good measure when performing bi-variant analysis

In [13]:

```
# Checking the share of Aspiration
```

```
sns.countplot(x = 'aspiration', data = df)  
plt.show()
```




There's an imbalance

In [14]:

```
# Checking the share of Door-number
```

```
sns.countplot(x = 'doornumber', data = df)  
plt.show()
```




More 4-door vehicles are available

In [15]:

```
# Checking the share of Car-body types
```

```
sns.countplot(x = 'carbody', data = df)  
plt.show()
```




Since 'hardtop' and 'convertible' are <5% of the data, they can be clubbed so that the model can also taken into account this minority class but need to check if the prices tend to foloow the same pattern

In [16]:

```
# Checking the share of Car body
```

```
sns.countplot(x = 'carbody', data = df)  
plt.show()
```



In [17]:

```
(df['carbody'].value_counts(normalize = True) * 100).round(2)
```

Out[17]:

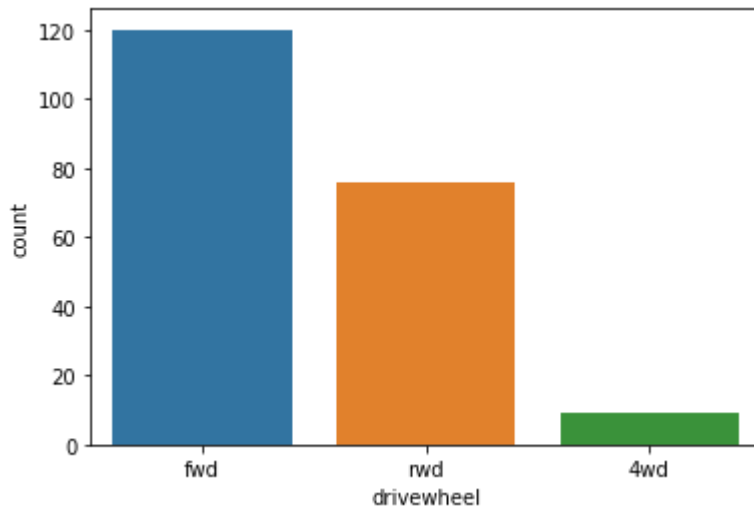
```
sedan          46.83  
hatchback      34.15  
wagon          12.20  
hardtop         3.90  
convertible     2.93  
Name: carbody, dtype: float64
```

'hardtop' and 'convertibles' need to be inspected for minority class pattern with the target variable

In [18]:

```
# Checking the distribution of Drive-wheel
```

```
sns.countplot(x = 'drivewheel', data = df, order = ['fwd', 'rwd', '4wd'])  
plt.show()
```




Forward Drives have more data points comparatively

In [19]:

```
# Checking the share of Engine Location
```

```
sns.countplot(x = 'enginelocation', data = df)  
plt.show()
```



In [20]:

```
(df['enginelocation'].value_counts(normalize = True) * 100).round(2)
```

Out[20]:

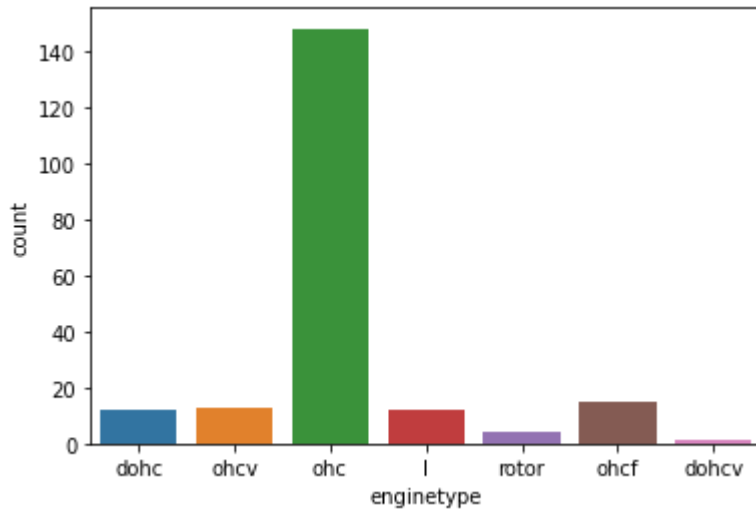
```
front    98.54  
rear      1.46  
Name: enginelocation, dtype: float64
```

High data imbalance, need to check for price for 'rear' type of vehicles if it's insightful then need to retain it else can be dropped

In [21]:

```
# Checking the share of Engine-type
```

```
sns.countplot(x = 'enginetype', data = df)  
plt.show()
```



In [22]:

```
(df['enginetype'].value_counts(normalize = True) * 100).round(2)
```

Out[22]:

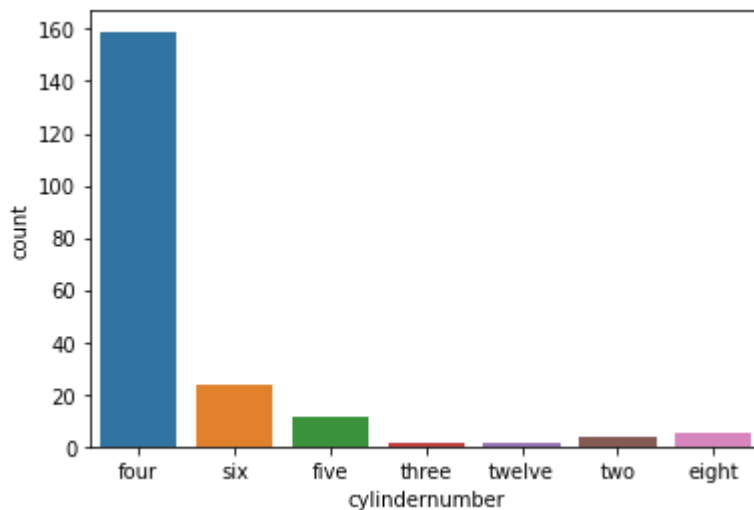
```
ohc      72.20  
ohcf     7.32  
ohcv     6.34  
l        5.85  
dohc     5.85  
rotor    1.95  
dohcv    0.49  
Name: enginetype, dtype: float64
```

'rotor' and 'dohcv' need to be inspected for minority class pattern with the target variable

In [23]:

```
# Checking the share of Cylinder-number
```

```
sns.countplot(x = 'cylindernumber', data = df)  
plt.show()
```



In [24]:

```
(df['cylindernumber'].value_counts(normalize = True) * 100).round(2)
```

Out[24]:


```
four      77.56  
six       11.71  
five       5.37  
eight      2.44  
two        1.95  
three      0.49  
twelve     0.49  
Name: cylindernumber, dtype: float64
```

'three', 'twelve', 'two' and 'eight' need to be inspected for minority class pattern with the target variable

In [25]:

```
# Checking the share of Fuel-system
```

```
sns.countplot(x = 'fuelsystem', data = df)  
plt.show()
```



In [26]:

```
(df['fuelsystem'].value_counts(normalize = True) * 100).round(2)
```

Out[26]:


```
mpfi    45.85  
2bbl    32.20  
idi      9.76  
1bbl     5.37  
spdi     4.39  
4bbl     1.46  
spfi     0.49  
mfi      0.49  
Name: fuelsystem, dtype: float64
```

'spfi', 'mfi', '4bbl' and 'spdi' need to be inspected for minority class pattern with the target variable

Bi-variant analysis

In [27]:

```
# Checking impact of Symboling (insurance risk factor) on median price  
sns.barplot(x = df['symboling'], y = df['price'], estimator = np.median)  
plt.show()
```



In [28]:

```
# Checking impact of Fuel type on median price
```

```
plt.figure(figsize=(15, 7))
```


```
plt.subplot(1, 2, 1)
```

```
sns.barplot(x = df['fueltype'], y = df['price'], estimator = np.median)
```

```
plt.subplot(1, 2, 2)
```

```
plt.pie(x = df['fueltype'].value_counts(), labels = df['fueltype'].value_counts.  
        shadow= True, explode = [0, 0.1], textprops= {'fontsize': 12, 'back  
        startangle = 45)
```

```
plt.show()
```



Diesel vehicles tend to cost higher though it's count is less


In [29]:

```
# Checking impact of Aspiration type on median price
```

```
plt.figure(figsize=(7, 7))
```

```
sns.barplot(x = df['aspiration'], y = df['price'], hue = df['fueltype'], estim
```

```
plt.show()
```



In [30]:

```
df['aspiration'].value_counts(normalize=True) * 100
```

Out[30]:

```
std      81.95122
turbo    18.04878
Name: aspiration, dtype: float64
```

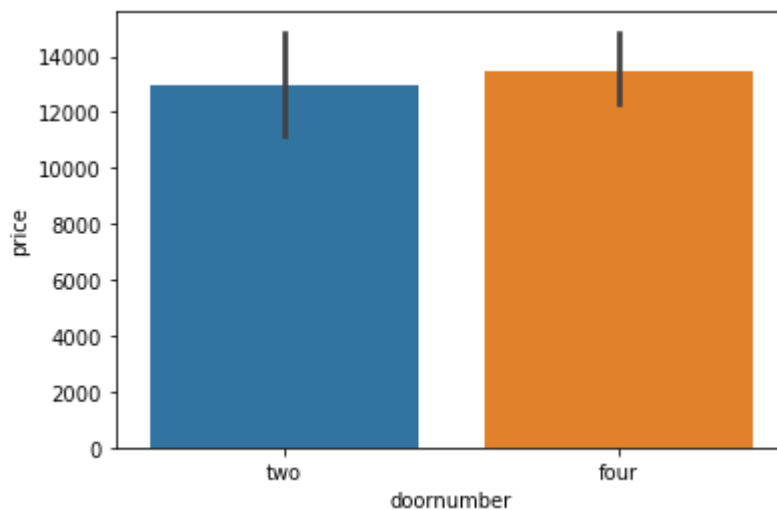
The difference in price between std and turbo for aspiration types is greater relatively due to which it's aggregated price is higher. Though the count of std is higher (by approx >4 times) yet its median price is low.

- Perhaps because std types are quite common

In [31]:

```
# Checking impact of Number of Doors on median price
```

```
sns.barplot(x = df['doornumber'], y = df['price'], estimator = np.mean)
plt.show()
```




The cost of both the types is approximately same

In [32]:

```
# Checking impact of Aspiration type on median price
```

```
sns.barplot(x = df['carbody'], y = df['price'], estimator = np.mean)
```

```
plt.show()
```




Since 'hardtop' and 'hatchback' don't follow the same pattern in prices, we can't combine them into a single category on the basis of minority classes

In [33]:

```
# Checking impact of Engine type on median price
```

```
sns.barplot(x = df['drivewheel'], y = df['price'], estimator = np.mean)
```

```
plt.show()
```



Foward and 4wd tend to follow similar pricing trend


Rear wheel drives has higher pricing

In [34]:

```
# Checking impact of Engine type on median price
```

```
sns.barplot(x = df['enginelocation'], y = df['price'], estimator = np.mean)
```

```
plt.show()
```




Though the count of rear type is in high minority level, yet it's average pricing is greater than twice the average of front types. Hence, this variable can prove as an important/significant variable for predicting the target

In [35]:

```
# Checking impact of Engine type on median price
```

```
sns.barplot(x = df['enginetype'], y = df['price'], estimator = np.mean)
```

```
plt.show()
```




In [36]:

```
# Checking impact of Number of cylinder on median price
```

```
sns.barplot(x = df['cylindernumber'], y = df['price'], estimator = np.mean)
```

```
plt.show()
```




Since 'three', 'twelve', 'two' and 'eight' don't follow the same pattern in prices, we can't combine them into a single category on the basis of minority classes

In [37]:

```
# Checking impact of Fuel type on median price
```

```
sns.barplot(x = df['fuelsystem'], y = df['price'], estimator = np.mean)
```

```
plt.show()
```



Since 'spfi', 'mfi', '4bbl' and 'spdi' follow the similar price pattern, they can be combined into a single category

In [38]:

```
(df['fuelsystem'].value_counts(normalize = True) * 100).round(2)
```

Out[38]:

```
mpfi    45.85
2bbl    32.20
idi      9.76
1bbl     5.37
spdi     4.39
4bbl     1.46
spfi     0.49
mfi      0.49
Name: fuelsystem, dtype: float64
```


Combining low % categories as others

In [39]:

```
df['fuelsystem'].replace(to_replace= ['spfi', 'mfi', '4bbl', 'spdi'],  
                        value= ['others', 'others', 'others', 'others'], inplace=True)
```

In [40]:

```
sns.countplot(x = 'fuelsystem', data = df)  
plt.show()
```



In [41]:

```
# Since 'symboling' should be treated as a categorical type ignoring it  
temp = df[df.columns[~(df.columns.isin(['symboling']))]]
```

In [42]:

```
temp.shape
```

Out[42]:

(205, 24)

In [43]:

```
temp.columns
```

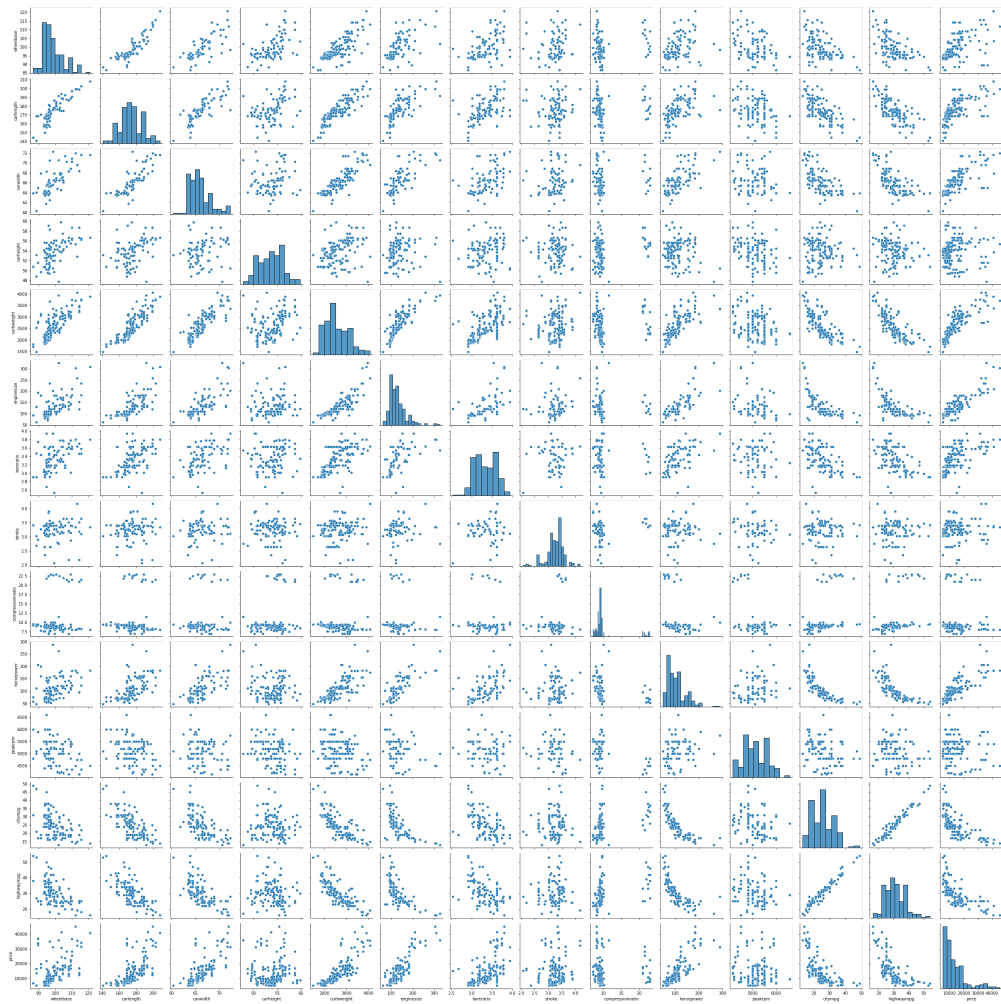
Out[43]:

```
Index(['make', 'fueltype', 'aspiration', 'doornumber', 'carbody', 'drivewheel',  
      'enginelocation', 'wheelbase', 'carlength', 'carwidth',  
      'carheight',  
      'curbweight', 'enginetype', 'cylindernumber', 'enginesize',  
      'fuelsystem', 'bore', 'stroke', 'compressionratio',  
      'horsepower',  
      'peakrpm', 'citympg', 'highwaympg', 'price'],  
      dtype='object')
```

In [44]:

```
sns.pairplot(temp)
```

```
plt.show()
```




In [45]:

```
plt.figure()

sns.pairplot(temp, y_vars = 'wheelbase', x_vars = ['carlength', 'curbweight', '
plt.show()
```

<Figure size 432x288 with 0 Axes>




In [46]:

```
plt.figure()

sns.pairplot(temp, y_vars = 'citympg', x_vars = ['horsepower', 'highwaympg'])

plt.show()
```

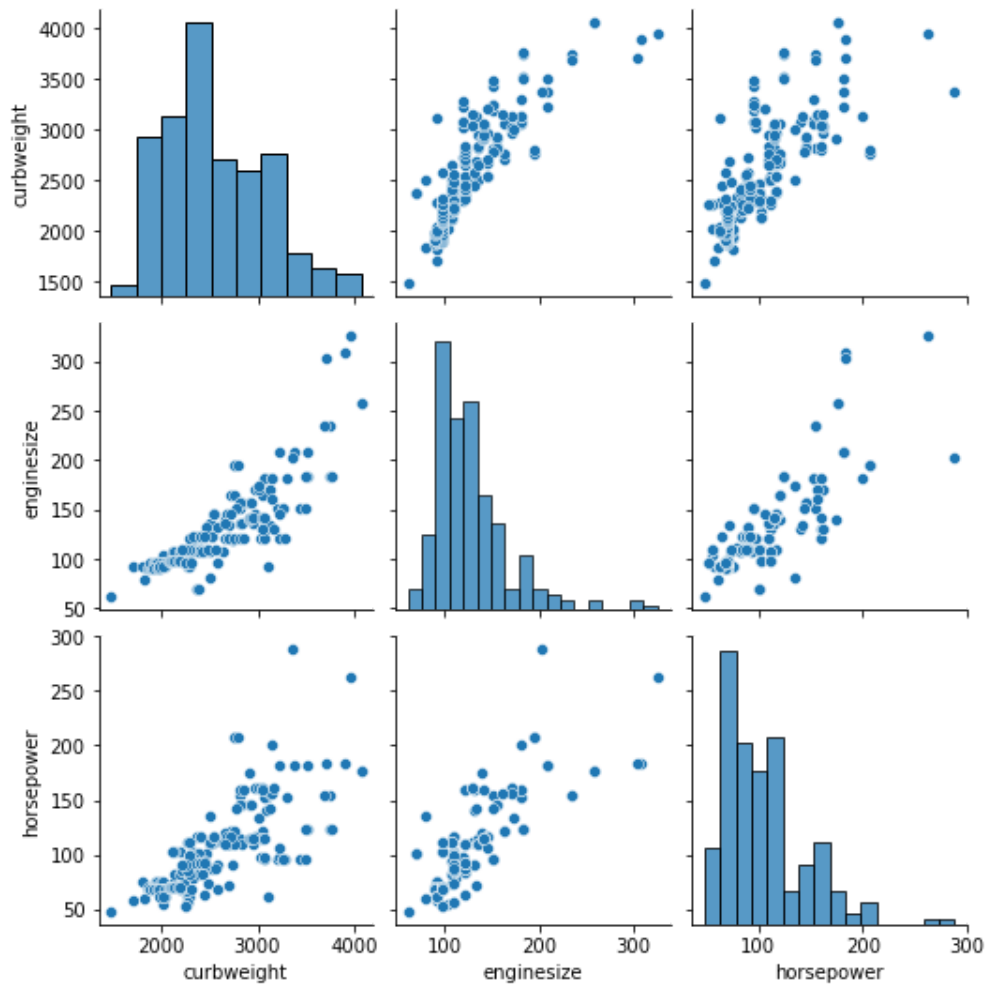
<Figure size 432x288 with 0 Axes>



citympg has a negative co-relation with horsepower


In [47]:

```
sns.pairplot(temp, vars = ['curbweight', 'enginesize', 'horsepower'])  
plt.show()
```



In [48]:

```
sns.pairplot(temp, y_vars = 'price', x_vars = ['curbweight', 'enginesize', 'horsepower', 'citympg', 'highwaympg'])  
plt.show()
```



curbweight, enginesize and horsepower follow sort of a linear co-relation with target variable with positive co-relation and seems to follow a similar trend

- Need to check for co-relation and VIF

citympg and highwaympg follow sort of a linear co-relation with target variable with negative co-relation and seems to follow a similar trend

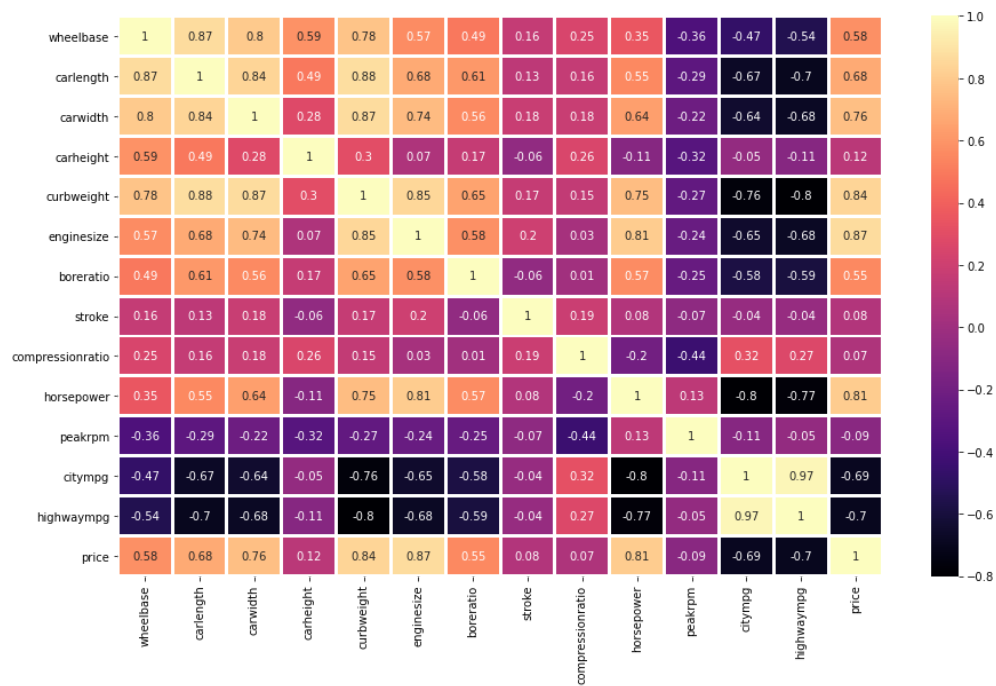
- Need to check for co-relation and VIF

In [49]:

```
# Co-relation Matrix
plt.figure(figsize=(15, 9))

sns.heatmap(temp.corr().round(2), cmap = 'magma', annot = True, linewidths = 2)

plt.show()
```



Wheelbase has high positive correlation with `carlength` , `carwidth` , `horsepower` and `boreratio` this also applies with each other and fairly good positive relation with `carheight`

Carlength has a mediumly positive co-relation with `enginesize` while mediumly negative with `citympg` and `highwaympg` . Applies to each other also

Enginesize has very high positive correlation with `horsepower` and `curbweight` (as expected)

Boreratio has quite high negative co-relation with `citympg` and `highwaympg`

CityMPG and HighwayMPG has high negative or weak positive co-relation with almost all variables

But very high positive co-relation with each

Price has very high positive co-relation with `enginesize` , `curbweight` , `horsepower` , `carwidth` , `carlength`

Price has very high negative co-relation with `highwaympg`

Feature Engineering

In [50]:

```
df['make'] = df['make'].apply(lambda x: x.split()[0])
```

In [51]:

```
df['make'].value_counts(ascending = True).index
```

Out[51]:

```
Index(['Nissan', 'toyouta', 'mercury', 'porcshce', 'vokswagen',  
      'maxda', 'vw',  
      'renault', 'alfa-romero', 'jaguar', 'chevrolet', 'isuzu',  
      'porsche',  
      'saab', 'plymouth', 'audi', 'bmw', 'buick', 'volkswagen',  
      'dodge',  
      'peugeot', 'volvo', 'subaru', 'honda', 'mitsubishi', 'maz  
da', 'nissan',  
      'toyota'],  
      dtype='object')
```

There's a few categories such as porsche and porcshce, toyouta and toyota, etc

- Both are the same, (spelling mistake/shortforms)

In [52]:

```
df['make'].nunique()
```

Out[52]:

28

In [53]:

```
df['make'].replace(to_replace = ['porcshce', 'toyouta', 'vw', 'vokswagen', 'max  
                             value = ['porsche', 'toyota', 'volkswagen', 'volkswagen', 'm
```

In [54]:


```
df['make'].nunique()
```

Out[54]:

22

In [55]:

```
plt.figure(figsize = (15, 7))  
  
sns.countplot(df['make'], order = df['make'].value_counts().index)  
  
plt.xticks(rotation = 45)  
plt.show()
```



Toyota has the highest count while mercury the lowest


Most manufactures lie between the range 6 - 12

In [56]:

```
pd.crosstab(df['make'], df['fueltype']).sort_values(by = ['diesel', 'gas']).plot
```

Out[56]:

<AxesSubplot:ylabel='make'>



Peugeot has almost equal number of cars in 'diesel' and 'petrol' types while also

being the highest in the 'diesel' category

Toyota has the highest count in 'petrol' category


There are no diesel vehicles post volvo

In [57]:

```
pd.crosstab(df['make'], df['enginelocation']).sort_values(by = ['rear', 'front'])
```

Out[57]:

<AxesSubplot:ylabel='make'>




Only Porsche has 'rear' engine type of cars

Outlier Analysis

In [58]:


```
sns.boxplot(x = 'symboling', data = df)  
plt.show()
```



No outliers present

In [59]:


```
sns.boxplot(x = 'wheelbase', data = df)  
plt.show()
```



Very few outliers

In [60]:


```
sns.boxplot(x = 'carlength', data = df)  
plt.show()
```



Very few outliers

In [61]:


```
sns.boxplot(x = 'carwidth', data = df)  
plt.show()
```



Very few outliers

In [62]:


```
sns.boxplot(x = 'carheight', data = df)  
plt.show()
```



No outliers present

In [63]:


```
sns.boxplot(x = 'curbweight', data = df)  
plt.show()
```



No outliers present

In [64]:


```
sns.boxplot(x = 'engine_size', data = df)  
plt.show()
```



Very few outliers

In [65]:


```
sns.boxplot(x = 'bore_ratio', data = df)  
plt.show()
```



No outliers present

In [66]:


```
sns.boxplot(x = 'stroke', data = df)  
plt.show()
```



Few outliers present, once it's co-relation is performed against bore-ratio. We can decide what to do

In [67]:

```
sns.boxplot(x = 'compressionratio', data = df)
plt.show()
```



Let's explore by restricting outlier

In [68]:

```
temp = df[df['compressionratio'] <= np.percentile(df['compressionratio'], 0.95)]
temp.shape
```

Out[68]:


(7, 25)

If we restrict to even 95th percentile, the number of records becomes lesser the one-fourth of the original data, and hence can't be dropped

Also, since majority of the data is above 95th percentile, we can't replace it with $1.5 \times \text{IQR}$ as it'll drastically increase the biasing

In [69]:


```
sns.boxplot(x = 'horsepower', data = df)  
plt.show()
```



Very few outliers

In [70]:


```
sns.boxplot(x = 'peakrpm', data = df)  
plt.show()
```



Very few outliers

In [71]:


```
sns.boxplot(x = 'citympg', data = df)  
plt.show()
```



Very few outliers

In [72]:


```
sns.boxplot(x = 'highwaympg', data = df)  
plt.show()
```



Very few outliers

In [73]:

```
sns.boxplot(x = 'price', data = df)
plt.show()
```



We have seen that, there's negligible outliers present in the data apart from a very few columns

Dummy Encoding

In [74]:

```
df['symboling'] = df['symboling'].astype('object')
```


In [75]:

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 205 entries, 0 to 204
Data columns (total 25 columns):
#   Column                Non-Null Count  Dtype
---  -
0   symboling              205 non-null    object
1   make                   205 non-null    object
2   fueltype               205 non-null    object
3   aspiration             205 non-null    object
4   doornumber             205 non-null    object
5   carbody                205 non-null    object
6   drivewheel            205 non-null    object
7   enginelocation         205 non-null    object
8   wheelbase              205 non-null    float64
9   carlength             205 non-null    float64
10  carwidth               205 non-null    float64
11  carheight              205 non-null    float64
12  curbweight             205 non-null    int64
13  enginetype             205 non-null    object
14  cylindernumber         205 non-null    object
15  enginesize             205 non-null    int64
16  fuelsystem             205 non-null    object
17  boreratio              205 non-null    float64
18  stroke                 205 non-null    float64
19  compressionratio       205 non-null    float64
20  horsepower             205 non-null    int64
21  peakrpm                205 non-null    int64
22  citympg                205 non-null    int64
23  highwaympg            205 non-null    int64
24  price                  205 non-null    float64
dtypes: float64(8), int64(6), object(11)
memory usage: 40.2+ KB
```


In [76]:

```
dum_df = pd.get_dummies(df.select_dtypes('object'), drop_first=True)
dum_df.head()
```

Out[76]:

	symboling_-1	symboling_0	symboling_1	symboling_2	symboling_3	make_a
0	0	0	0	0	1	
1	0	0	0	0	1	
2	0	0	1	0	0	
3	0	0	0	1	0	
4	0	0	0	1	0	

5 rows × 52 columns




In [77]:

```
df = pd.concat([df.loc[:, ~df.columns.isin(df.select_dtypes('object').columns)]
df.head()
```

Out[77]:

	wheelbase	carlength	carwidth	carheight	curbweight	enginesize	boreratio
0	88.6	168.8	64.1	48.8	2548	130	3.47
1	88.6	168.8	64.1	48.8	2548	130	3.47
2	94.5	171.2	65.5	52.4	2823	152	2.68
3	99.8	176.6	66.2	54.3	2337	109	3.19
4	99.4	176.6	66.4	54.3	2824	136	3.19

5 rows × 66 columns



Dividing train and test data

In [78]:

```
# Splitting the data
```

```
train, test = train_test_split(df, test_size = 0.3, random_state = 100)
```

In [79]:

```
df.shape
```

Out[79]:

```
(205, 66)
```

In [80]:

```
train.shape
```

Out[80]:

```
(143, 66)
```

In [81]:

```
test.shape
```

Out[81]:

```
(62, 66)
```

Normalising the data

In [82]:

```
norm_vars = ['wheelbase', 'carlength', 'carwidth', 'carheight', 'curbweight', 'stroke', 'compressionratio', 'horsepower', 'peakrpm', 'citympg']
```

Since most of the numerical independent variables are quite normally distributed and there are no columns where the values are strictly between 0 and 1 (except dummy variables), we can go ahead and apply standardization

In [83]:


```
scaler = StandardScaler()

train[norm_vars] = scaler.fit_transform(train[norm_vars])
train.head()
```

Out[83]:

	wheelbase	carlength	carwidth	carheight	curbweight	enginesize	borera
122	-0.811836	-0.487238	-0.924500	-1.134628	-0.642128	-0.660242	-1.2973
125	-0.677177	-0.359789	1.114978	-1.382026	0.439415	0.637806	2.4322
166	-0.677177	-0.375720	-0.833856	-0.392434	-0.441296	-0.660242	-0.2591
1	-1.670284	-0.367754	-0.788535	-1.959288	0.015642	0.123485	0.6251
199	0.972390	1.225364	0.616439	1.627983	1.137720	0.123485	1.2018

5 rows × 66 columns




In [84]:

```
train.describe().round(2)
```

Out[84]:

	wheelbase	carlength	carwidth	carheight	curbweight	enginesize	borera
count	143.00	143.00	143.00	143.00	143.00	143.00	143.00
mean	0.00	0.00	-0.00	0.00	-0.00	-0.00	-0.00
std	1.00	1.00	1.00	1.00	1.00	1.00	1.00
min	-2.01	-2.57	-2.51	-2.37	-1.94	-1.57	-2.00
25%	-0.68	-0.62	-0.86	-0.72	-0.77	-0.68	-0.68
50%	-0.34	-0.11	-0.20	0.06	-0.25	-0.37	0.00
75%	0.45	0.71	0.48	0.74	0.72	0.39	0.00
max	2.87	2.32	2.93	2.29	2.81	4.92	2.00

8 rows × 66 columns



In [85]:

```
test[norm_vars] = scaler.transform(test[norm_vars])
```

In [86]:

```
# Dividing the predictors and target variable for train dataset
```

```
y_train = train.pop('price')  
X_train = train
```

```
print(y_train.shape)  
print(X_train.shape)
```

```
(143,)  
(143, 65)
```

In [87]:

```
# Dividing the predictors and target variable for test dataset
```

```
y_test = test.pop('price')  
X_test = test  
print(y_test.shape)  
print(X_test.shape)
```

```
(62,)  
(62, 65)
```

6. Model Building

Now, building model by using Recursive Feature Engineering (RFE), selecting top 15 variables that are significant.

In [88]:

```
# Running RFE with the output number of the variable equal to 6
```

```
lm = LinearRegression()  
lm.fit(X_train, y_train)
```

```
rfe = RFE(lm, 15)  
rfe = rfe.fit(X_train, y_train)
```

In [89]:

```
# Overview of significant variables
```

```
list(zip(X_train.columns,rfe.support_,rfe.ranking_))
```

Out[89]:

```
[('wheelbase', False, 27),
 ('carlength', False, 22),
 ('carwidth', False, 12),
 ('carheight', False, 25),
 ('curbweight', False, 18),
 ('enginesize', True, 1),
 ('boreratio', False, 9),
 ('stroke', False, 17),
 ('compressionratio', False, 33),
 ('horsepower', False, 41),
 ('peakrpm', False, 35),
 ('citympg', False, 48),
 ('highwaympg', False, 42),
 ('symboling_1', True, 1),
 ('symboling_0', True, 1),
 ('symboling_1', True, 1),
 ('symboling_2', True, 1),
 ('symboling_3', True, 1),
 ('make_audi', True, 1),
 ('make_bmw', True, 1),
 ('make_buick', True, 1),
 ('make_chevrolet', False, 24),
 ('make_dodge', False, 20),
 ('make_honda', False, 23),
 ('make_isuzu', False, 45),
 ('make_jaguar', False, 26),
 ('make_mazda', False, 39),
 ('make_mercury', False, 51),
 ('make_mitsubishi', False, 11),
 ('make_nissan', False, 36),
 ('make_peugeot', False, 10),
 ('make_plymouth', False, 19),
 ('make_porsche', True, 1),
 ('make_renault', False, 49),
 ('make_saab', True, 1),
 ('make_subaru', False, 14),
 ('make_toyota', False, 40),
 ('make_volkswagen', False, 38),
 ('make_volvo', True, 1),
 ('fueltype_gas', False, 31),
 ('aspiration_turbo', False, 15),
 ('doornumber_two', False, 44),
```

```
('carbody_hardtop', False, 28),
('carbody_hatchback', False, 21),
('carbody_sedan', False, 29),
('carbody_wagon', False, 30),
('drivewheel_fwd', False, 50),
('drivewheel_rwd', False, 37),
('enginelocation_rear', True, 1),
('enginetype_dohcv', False, 4),
('enginetype_l', False, 7),
('enginetype_ohc', False, 43),
('enginetype_ohcf', False, 8),
('enginetype_ohcv', False, 13),
('enginetype_rotor', True, 1),
('cylindernumber_five', False, 6),
('cylindernumber_four', False, 5),
('cylindernumber_six', False, 16),
('cylindernumber_three', False, 3),
('cylindernumber_twelve', False, 2),
('cylindernumber_two', True, 1),
('fuelsystem_2bbl', False, 34),
('fuelsystem_idi', False, 32),
('fuelsystem_mpf', False, 47),
('fuelsystem_others', False, 46)]
```

Fine Tuning Model

- **Checking Multicollinearity among significant variables**

In [90]:

```
X_train_rfe = X_train[X_train.columns[rfe.support_]]
```

In [91]:

```
vif = pd.DataFrame()
vif['Features'] = X_train_rfe.columns
vif['VIF'] = [variance_inflation_factor(X_train_rfe.values, i) for i in range(X_train_rfe.shape[1])]
vif['VIF'] = round(vif['VIF'], 2)
vif = vif.sort_values(by = "VIF", ascending = False)
vif
```

Out[91]:

	Features	VIF
13	enginetype_rotor	inf
14	cylindernumber_two	inf
1	symboling_-1	1.84
0	enginesize	1.68
8	make_buick	1.65
9	make_porsche	1.65
5	symboling_3	1.54
12	enginelocation_rear	1.53
11	make_volvo	1.39
4	symboling_2	1.21
3	symboling_1	1.19
7	make_bmw	1.18
10	make_saab	1.13
6	make_audi	1.10
2	symboling_0	1.07

In [92]:

```
# Intercept addition
X_train_rfe = sm.add_constant(X_train_rfe)
```


In [93]:

```
# Training the model
```

```
lm1 = sm.OLS(y_train,X_train_rfe).fit()
```

In [94]:

```
# Viewing parameter's co-efficients and significance
```

```
lm1.summary()
```

Out[94]:

OLS Regression Results

Dep. Variable:	price	R-squared:	0.916			
Model:	OLS	Adj. R-squared:	0.906			
Method:	Least Squares	F-statistic:	99.29			
Date:	Sun, 24 Jan 2021	Prob (F-statistic):	1.42e-61			
Time:	17:55:20	Log-Likelihood:	-26.081			
No. Observations:	143	AIC:	82.16			
Df Residuals:	128	BIC:	126.6			
Df Model:	14					
Covariance Type:	nonrobust					
	coef	std err	t	P> t 	[0.025	0.975]
const	-1.0705	0.295	-3.628	0.000	-1.654	-0.487
enginesize	0.7018	0.033	21.101	0.000	0.636	0.768
symboling_-1	0.7821	0.266	2.942	0.004	0.256	1.308
symboling_0	0.9325	0.298	3.125	0.002	0.342	1.523
symboling_1	0.8277	0.300	2.763	0.007	0.235	1.420
symboling_2	0.7018	0.306	2.293	0.024	0.096	1.308
symboling_3	0.8325	0.305	2.731	0.007	0.229	1.436
make_audi	0.8073	0.144	5.595	0.000	0.522	1.093
make_bmw	1.2020	0.136	8.835	0.000	0.933	1.471
make_buick	1.2273	0.182	6.745	0.000	0.867	1.587
make_porsche	1.1008	0.228	4.835	0.000	0.650	1.551
make_saab	0.6850	0.188	3.644	0.000	0.313	1.057
make_volvo	0.9755	0.200	4.874	0.000	0.579	1.372
enginelocation_rear	0.6462	0.379	1.704	0.091	-0.104	1.397
enginetype_rotor	0.5675	0.093	6.119	0.000	0.384	0.751
cylindernumber_two	0.5675	0.093	6.119	0.000	0.384	0.751

Omnibus:	16.056	Durbin-Watson:	1.973
Prob(Omnibus):	0.000	Jarque-Bera (JB):	27.045
Skew:	0.544	Prob(JB):	1.34e-06
Kurtosis:	4.832	Cond. No.	3.67e+16

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

[2] The smallest eigenvalue is 1.35e-31. This might indicate that there are strong multicollinearity problems or that the design matrix is singular.

In [95]:

```
X_train_rfe.shape
```

Out[95]:

```
(143, 16)
```

In [96]:

```
# Dropping 'enginetype_dohcv'  
X_train_rfe.drop('engine_location_rear', axis = 1, inplace = True)
```

In [97]:

```
X_train_rfe.shape
```

Out[97]:

```
(143, 15)
```

In [98]:

```
# Adding constant, training model and viewing parameter's co-efficients and sig  
X_train_rfe = sm.add_constant(X_train_rfe)  
  
lm1a = sm.OLS(y_train,X_train_rfe).fit()  
  
lm1a.summary()
```

Out[98]:

OLS Regression Results

Dep. Variable:	price	R-squared:	0.914
Model:	OLS	Adj. R-squared:	0.905
Method:	Least Squares	F-statistic:	105.1
Date:	Sun, 24 Jan 2021	Prob (F-statistic):	5.62e-62
Time:	17:55:20	Log-Likelihood:	-27.684
No. Observations:	143	AIC:	83.37
Df Residuals:	129	BIC:	124.8
Df Model:	13		
Covariance Type:	nonrobust		

	coef	std err	t	P> t	[0.025	0.975]
const	-1.0681	0.297	-3.594	0.000	-1.656	-0.480
enginesize	0.7017	0.034	20.944	0.000	0.635	0.768
symboling_-1	0.7821	0.268	2.920	0.004	0.252	1.312
symboling_0	0.9299	0.301	3.094	0.002	0.335	1.525
symboling_1	0.8204	0.302	2.719	0.007	0.223	1.417
symboling_2	0.6998	0.308	2.269	0.025	0.090	1.310
symboling_3	0.8478	0.307	2.762	0.007	0.241	1.455
make_audi	0.8092	0.145	5.567	0.000	0.522	1.097
make_bmw	1.2030	0.137	8.777	0.000	0.932	1.474
make_buick	1.2220	0.183	6.668	0.000	0.859	1.585
make_porsche	1.3062	0.195	6.714	0.000	0.921	1.691
make_saab	0.6788	0.189	3.586	0.000	0.304	1.053
make_volvo	0.9731	0.202	4.827	0.000	0.574	1.372

enginetype_rotor	0.5586	0.093	5.988	0.000	0.374	0.743
cylindernumber_two	0.5586	0.093	5.988	0.000	0.374	0.743

Omnibus:	14.688	Durbin-Watson:	2.006
Prob(Omnibus):	0.001	Jarque-Bera (JB):	23.096
Skew:	0.524	Prob(JB):	9.66e-06
Kurtosis:	4.667	Cond. No.	1.34e+17

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

[2] The smallest eigenvalue is 1.01e-32. This might indicate that there are strong multicollinearity problems or that the design matrix is singular.

In [99]:

```
X_train_rfe.drop('const', axis = 1, inplace = True)
vif = pd.DataFrame()
vif['Features'] = X_train_rfe.columns
vif['VIF'] = [variance_inflation_factor(X_train_rfe.values, i) for i in range(X_train_rfe.shape[0])]
vif['VIF'] = round(vif['VIF'], 2)
vif = vif.sort_values(by = "VIF", ascending = False)
vif
```

Out[99]:

	Features	VIF
12	enginetype_rotor	inf
13	cylindernumber_two	inf
1	symboling_-1	1.84
0	enginesize	1.68
8	make_buick	1.65
5	symboling_3	1.52
11	make_volvo	1.39
4	symboling_2	1.21
3	symboling_1	1.19
9	make_porsche	1.19
7	make_bmw	1.18
10	make_saab	1.12
6	make_audi	1.10
2	symboling_0	1.07

In [100]:

```
X_train_rfe.shape
```

Out[100]:

(143, 14)

In [101]:

```
# Dropping 'cylindernumber_five'
X_train_rfe.drop('enginetype_rotor', axis = 1, inplace = True)
```

In [102]:

```
X_train_rfe.shape
```

Out[102]:

```
(143, 13)
```

In [103]:

```
# Adding constant, training model and viewing parameter's co-efficients and sig  
X_train_rfe = sm.add_constant(X_train_rfe)  
  
lm1b = sm.OLS(y_train,X_train_rfe).fit()  
  
lm1b.summary()
```

Out[103]:

OLS Regression Results

Dep. Variable:	price	R-squared:	0.914
Model:	OLS	Adj. R-squared:	0.905
Method:	Least Squares	F-statistic:	105.1
Date:	Sun, 24 Jan 2021	Prob (F-statistic):	5.62e-62
Time:	17:55:21	Log-Likelihood:	-27.684
No. Observations:	143	AIC:	83.37
Df Residuals:	129	BIC:	124.8
Df Model:	13		
Covariance Type:	nonrobust		

	coef	std err	t	P> t	[0.025	0.975]
const	-1.0681	0.297	-3.594	0.000	-1.656	-0.480
enginesize	0.7017	0.034	20.944	0.000	0.635	0.768
symboling_-1	0.7821	0.268	2.920	0.004	0.252	1.312
symboling_0	0.9299	0.301	3.094	0.002	0.335	1.525
symboling_1	0.8204	0.302	2.719	0.007	0.223	1.417
symboling_2	0.6998	0.308	2.269	0.025	0.090	1.310
symboling_3	0.8478	0.307	2.762	0.007	0.241	1.455
make_audi	0.8092	0.145	5.567	0.000	0.522	1.097
make_bmw	1.2030	0.137	8.777	0.000	0.932	1.474
make_buick	1.2220	0.183	6.668	0.000	0.859	1.585
make_porsche	1.3062	0.195	6.714	0.000	0.921	1.691
make_saab	0.6788	0.189	3.586	0.000	0.304	1.053
make_volvo	0.9731	0.202	4.827	0.000	0.574	1.372

cylindernumber_two 1.1171 0.187 5.988 0.000 0.748 1.486

Omnibus: 14.688 Durbin-Watson: 2.006
Prob(Omnibus): 0.001 Jarque-Bera (JB): 23.096
Skew: 0.524 Prob(JB): 9.66e-06
Kurtosis: 4.667 Cond. No. 31.5

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

In [104]:

```
X_train_rfe.drop('const', axis = 1, inplace = True)
vif = pd.DataFrame()
vif['Features'] = X_train_rfe.columns
vif['VIF'] = [variance_inflation_factor(X_train_rfe.values, i) for i in range(X_train_rfe.shape[1])]
vif['VIF'] = round(vif['VIF'], 2)
vif = vif.sort_values(by = "VIF", ascending = False)
vif
```

Out[104]:

	Features	VIF
1	symboling_-1	1.84
0	enginesize	1.68
8	make_buick	1.65
5	symboling_3	1.52
12	cylindernumber_two	1.46
11	make_volvo	1.39
4	symboling_2	1.21
3	symboling_1	1.19
9	make_porsche	1.19
7	make_bmw	1.18
10	make_saab	1.12
6	make_audi	1.10
2	symboling_0	1.07

In [105]:

```
X_train_rfe.shape
```

Out[105]:

```
(143, 13)
```

model lm1g with variables symboling_-1 , enginesize , make_buick , symboling_3 , cylindernumber_two , make_volvo , symboling_2 , symboling_1 , make_porsche , make_bmw , make_saab , make_audi , symboling_0 have VIF < 2 and all variables are significant The R squared and adjusted R squared are also almost same. This is a good model. Hence proceeding with residual analysis.

*** Equation *:**

$$x = -1.068 + \text{enginesize} \times 0.7017 + \text{symboling_} \mathbf{-1} \times 0.7821 + \text{symboling_} \mathbf{0} \times 0.6998 + \text{symboling_} \mathbf{3} \times 0.8478 + \text{make_audi} \times 0.8092 + \text{make_bmw} \times 1.203 + 1.3062 + \text{make_saab} \times 0.6788 + \text{make_volvo} \times 0.9731 + \text{cylindernumber_tw}$$

In [106]:

```
# Viewing Predictor's co-efficients
```

```
lm1b.params
```

Out[106]:

```
const          -1.068057
enginesize      0.701693
symboling_-1    0.782086
symboling_0     0.929943
symboling_1     0.820362
symboling_2     0.699815
symboling_3     0.847760
make_audi       0.809185
make_bmw        1.202960
make_buick      1.221990
make_porsche    1.306203
make_saab       0.678774
make_volvo      0.973115
cylindernumber_two 1.117143
dtype: float64
```

a. Residual Analysis - Train Data

In [107]:

```
# Predicting using train data
```

```
X_train_rfe = sm.add_constant(X_train_rfe)
```

```
y_train_pred = lm1b.predict(X_train_rfe)
```

In [108]:

```
# Residual calculation
```

```
residuals = y_train - y_train_pred
```

In [109]:

```
#Checking if Residuals are normally distributed
```


```
plt.figure(figsize = [6,6])
```

```
sns.distplot(residuals, bins = 7)
```

```
plt.title('Probabilty Distribution of Residuals - Train data')
```

```
plt.ylabel('')
```

```
plt.show()
```



In [110]:


```
# Checking residual are randomly distributed (no pattern)

plt.figure(figsize = [6,6])

sns.regplot(y_train, y_train_pred, line_kws = {'color': 'r'})

plt.title('Train Data', fontsize = 18)
plt.xlabel('Predicted Y', fontsize = 15)
plt.ylabel('Actual Y', fontsize = 15)

plt.show()
```



From the above plots we can infer that the error terms are normally distributed about the mean ~ 0 and the predicted and actual values have a almost linear relation hence the model is a good fit and predicts data well.


In [111]:

```
# Checking for residual normality using Q-Q plot on test data
```

```
sm.qqplot(residuals, fit = True, line = '45')
```

```
plt.title('Train Data')  
plt.xlabel('Residuals')  
plt.ylabel('Normal data')
```

```
plt.show()
```




In [112]:

```
# Checking for homoscedasticity on test data

sns.regplot(y_train_pred, residuals, line_kws = {'color': 'r'})

plt.title('Train Data')
plt.xlabel('Fitted Values')
plt.ylabel('Residuals')

plt.show()
```



Evaluation on Test data

In [113]:

```
# Columns required for model prediction

req_cols = vif['Features']
```

In [114]:

```
# Constant addition and Predicting

X_test=sm.add_constant(X_test[req_cols])
y_test_pred=lm1b.predict(X_test)
```

b. Residual Analysis - Test Data

In [115]:

```
residuals = y_test - y_test_pred
```

In [116]:

```
# Checking if Residuals are normally distributed
```


```
plt.figure(figsize = [6,6])
```

```
sns.distplot(residuals)
```

```
plt.title('Probabilty Distribution of Residuals - Test Data')
```

```
plt.ylabel('')
```

```
plt.show()
```



In [117]:


```
# Checking residual are randomly distributed (no pattern)

plt.figure(figsize = [6,6])

sns.regplot(y_test, y_test_pred, line_kws = {'color': 'r'})

plt.title('Test Data', fontsize = 18)
plt.xlabel('Predicted Y', fontsize = 15)
plt.ylabel('Actual Y', fontsize = 15)

plt.show()
```




In [118]:

```
# Checking for residual normality using Q-Q plot on test data

sm.qqplot(residuals, fit = True, line = '45')

plt.title('Test Data')
plt.xlabel('Residuals')
plt.ylabel('Normal data')

plt.show()
```



Comparing R^2 and Adjusted R^2 between Train and Test data

Testing for over-fitting

In [119]:

```
# R2 - train dataset

print(r2_score(y_train, y_train_pred).round(4) * 100, '%')
```

91.38 %

In [120]:

```
# R2 - test dataset  
print(r2_score(y_test, y_test_pred).round(4) * 100, '%')
```

85.65 %

$$\text{Adjusted } R^2 = 1 - \frac{(1 - R^2)(N - 1)}{N - p - 1}$$

In [121]:

```
# Adjusted R2 - train dataset  
print((1 - ((1-r2_score(y_train, y_train_pred))*(X_train_rfe.shape[0]))/(X_train  
X_train
```

90.36999999999999 %

In [122]:

```
# Adjusted R2 - test dataset  
print((1 - ((1-r2_score(y_train, y_train_pred))*(X_test.shape[0]))/(X_test.shape  
X_test.
```

88.62 %

Measurement	Train Dataset	Test Dataset
R^2	91.57 %	85.65 %
Adjusted R^2	90.37 %	88.62 %

As seen above the model does a little better on generalisation on the train data.

